

Sparse-Concept Calibration of Knowledge Tracing Models for Threshold-Based Educational Decisions

Khanh-Trinh Nguyen¹, Tuan Dao Minh¹, Duong Nguyen Tien¹, Chi Thanh Nguyen², and Van-Hau Nguyen^{1,*}

¹ Hung Yen University of Technology and Education, Hung Yen, Vietnam

² Academy of Military Science and Technology, Ha Noi, Vietnam

Email: trinhnk@utehy.edu.vn (K.-T.N.); tuanymc@utehy.edu.vn (T.D.M.); duongnt@utehy.edu.vn (D.N.T.); thanhnc@ioit.ai.vn (C.T.N.); haunv@utehy.edu.vn (V.-H.N.)

*Corresponding author

Abstract—Knowledge Tracing (KT) systems that skip, remediate, or advance practice typically consume a predicted probability rather than an area-under-the-curve (AUC) score, so population ranking can look acceptable while probabilities are poorly calibrated on rarely practiced knowledge components (KCs). This paper reports a diagnostic evaluation—not a new KT architecture—of train-only KC-frequency strata on three public logs (ASSISTments 2012, Junyi Academy, and XES3G5M) with IRT, DKT, and SimpleKT. Lower KC training frequency does not universally degrade discrimination, but calibration can become less reliable in some sparse-concept regimes. On ASSISTments 2012, SimpleKT expected calibration error (ECE) rises from 0.114 on dense KCs to 0.228 on sparse KCs (Limited occupancy, $N \approx 415$); Junyi has no learner-based sparse stratum, and XES3G5M SimpleKT ECE is essentially flat. A locked global threshold at $\tau=0.7$ raises the SimpleKT false-advance rate (incorrect-response rate among advance decisions) from 0.196 (dense) to 0.268 (sparse) on one ASSISTments fold; the sparse–dense gap stays positive on all five training runs across four unique learner partitions (mean 0.047). This is a simulated decision gate, not a classroom intervention. Probability-threshold decisions should be validated by KC-frequency stratum rather than on population AUC or ECE alone.

Keywords—knowledge tracing, calibration, sparse concepts, learning analytics, educational decision support, mastery threshold

I. INTRODUCTION

Adaptive practice platforms and intelligent tutoring systems use Knowledge Tracing (KT) to estimate a learner’s command of a skill and to decide whether to skip, remediate, or advance practice [1]–[3]. Those systems rarely consume a population area-under-the-curve (AUC) statistic. They consume a predicted probability p , which is then compared with a threshold τ . If $p \geq \tau$, the platform typically withholds additional practice on that skill. The operational chain is therefore population AUC $\rightarrow$ predicted probability $\rightarrow$ calibration $\rightarrow$ threshold-based educational decisions, with particular diagnostic risk on sparsely trained knowledge components (KCs).

KT evaluations typically report population AUC and accuracy [4], [5]. Those metrics rank models, but they can conceal miscalibration on low-frequency KCs—skills with little training-fold evidence—because most test events lie on dense, frequently practiced concepts. A model can discriminate well in aggregate and still assign overconfident probabilities on the sparse tail. When a fixed threshold is applied to p , that overconfidence appears as an advance decision followed by an incorrect next response. We term this the false-advance rate (FAR). Section III defines FAR as $P(y=0 \mid p \geq \tau)$; y denotes observed next-response correctness, not latent mastery truth.

This paper treats that mismatch as an evaluation problem for educational technology, not as a reason to propose another KT architecture. The study is organized around three research questions. RQ1: Does lower KC training frequency systematically degrade predictive discrimination? RQ2: How does calibration vary across KC-frequency strata and datasets? RQ3: When a fixed probability threshold is applied, does decision-error behavior differ between sparse and dense KCs?

The empirical answers are bounded. Lower KC training frequency does not universally degrade discrimination: on XES3G5M, sparse AUC is higher than dense AUC for DKT and SimpleKT. Calibration can, however, become less reliable in some sparse-concept regimes. On ASSISTments 2012, SimpleKT expected calibration error (ECE) increases from dense to sparse KCs, and a locked gate at $\tau=0.7$ yields a higher FAR on sparse than on dense advances. The same ECE gradient is absent on Junyi, where the learner-based sparse stratum is empty, and is essentially absent for SimpleKT on XES3G5M. We therefore do not claim that sparse KCs always fail, and we do not claim a causal effect of training frequency on calibration.

Contributions are conservative: (i) a train-only KC-frequency protocol with an explicit strict cold-start group, so the definition of “sparse” cannot leak test-fold counts; (ii) per-stratum calibration (ECE and Brier decomposition) on three public datasets, with occupancy flags (Reliable / Limited / Insufficient); (iii) a locked-threshold simulation of FAR and miss rates, with a five-seed check of the sparse–dense FAR gap on ASSISTments 2012. We do not propose a new KT architecture, a new calibration algorithm, or a new auditing theory, and we do not report a classroom intervention. A graph-KT (GKT) model and a CL4KT-style adapter appear only as an exploratory single-fold diagnostic on ASSISTments.

II. LITERATURE REVIEW

A. Knowledge Tracing and benchmark models

Knowledge Tracing (KT) models a learner’s evolving proficiency from historical interactions in order to predict the next response [3]. Classical approaches include Bayesian Knowledge Tracing (BKT), which tracks a latent mastery state for each skill [1], and item-response theory (IRT), of which the one-parameter Rasch model is the canonical form used in this study [6]. Deep Knowledge Tracing (DKT) replaced hand-specified mastery dynamics with a recurrent sequence model and established the modern next-response prediction task [2]. Attention-based successors, including context-aware attentive KT [7] and SimpleKT [4], retain that

task while substituting self-attention for recurrence. Benchmarking libraries such as pyKT have standardized preprocessing and comparison, and have shown that evaluation choices can distort reported discrimination [5]. Across this line of work, the primary reported statistic remains population AUC.

B. Graph and self-supervised KT

Graph-based KT (GKT) represents knowledge components as nodes of a graph and updates proficiency with a graph neural network [8]. Contrastive learning for KT (CL4KT) trains representations from augmented views of learning histories [9]. These models are related architectures, not the contribution of the present paper. They are included later only as an exploratory single-fold diagnostic on ASSISTments 2012, not as a proposed method and not as a state-of-the-art comparison.

C. Sparse-data and cold-start problems

The term “sparse” is used in several incompatible senses in the KT literature. This paper distinguishes four. (1) Sparse attention: sparseKT applies k-sparse attention so that only a subset of historical interactions receives nonzero weight [16]. (2) Sparse KC frequency: a knowledge component with few training-fold observations. (3) New-student cold start: prediction for learners who have little or no interaction history at test time [17]. (4) Concept-level cold start: a knowledge component with zero training-fold frequency, regardless of whether the learner is new. Educational interaction logs additionally exhibit a long tail in which a small set of dense KCs dominates event volume [20].

These four problems are not interchangeable. sparseKT-style sparse attention is not equivalent to low-frequency KCs: attention can be sparsified on a frequently practiced skill. New-student cold start is not the same as concept-level zero-frequency cold start: a previously observed learner may still encounter a concept that never appeared in the training fold. The protocol in Section III operationalizes (2) and (4) through train-only KC-frequency strata with an explicit strict cold-start group ($f=0$).

D. Calibration and educational decision support

When predicted probabilities are consumed as educational decisions, ranking is not a sufficient evaluation [10]. Calibration asks whether the predicted probability p matches the empirical frequency of a correct next response. Expected calibration error (ECE) is the usual bin-wise summary of that mismatch [12]; modern neural networks can remain poorly calibrated even when they discriminate well [11]. The Brier score [13] and its reliability, resolution, and uncertainty decomposition [14] separate miscalibration from task difficulty, and reliability diagrams display the same probability-level comparison. Complementary post-hoc work has begun to correct per-item bias after a frozen KT backbone [15]. Probability-level evaluation of this kind is still typically reported on the pooled population, not on train-only frequency strata.

Existing KT benchmarks mainly emphasize aggregate discrimination. Sparse frequency, calibration, sample support, and threshold-based decision error are rarely examined together. This study keeps standard models and reports those four quantities jointly, so that a population AUC

result can be read as an educational-technology decision rather than only as a leaderboard rank.

III. MATERIALS AND METHODS

A. Datasets

We use three de-identified public logs after a common interaction schema: ASSISTments 2012 [18], [5], Junyi Academy [19], and XES3G5M [20]. Table 1 reports post-processing counts: processed learners, processed KCs, processed interactions, and learner-based test events. Exact processed interaction totals are 2,657,490 (ASSISTments 2012), 16,215,567 (Junyi Academy), and 7,953,709 (XES3G5M). Preprocessing drops rows with missing learner, KC, or label identifiers, binarizes correctness, parses timestamps, and keeps learners with at least two interactions. Those filters—not the raw public dumps—are what Table 1 counts.

Provenance. ASSISTments 2012 is the 2012–2013 public release with affect [18], documented for KT use in pyKT [5]; each retained row already carries one `skill_id`. Junyi Academy is the Kaggle Online Learning Activity Dataset [19] (not Junyi2015); the KC field is `ucid`. XES3G5M [20] is loaded from the authors’ `kc_level` split: `train_valid_sequences.csv` and `test.csv` are concatenated and flattened to one row per sequence position.

XES3G5M counts. The original dataset paper reports 18,066 students, 7,652 questions, 865 KCs, and 5,549,635 question-level interactions [20]. Table 1 instead lists 866 KCs and 7.95M interactions because flattening `kc_level` sequences (i) expands multi-KC questions into one row per listed concept and (ii) retains sequence-padding tokens coded `skill_id=-1` (with matching `question_id=-1`). Unique `kc_id` including that padding token is 866; excluding it is 865. Unique `item_id` including `-1` is 7,653. Subsequent preprocessing dropped 0 rows. We do not retabulate Table 1 to the question-level official totals.

TABLE 1. POST-PROCESSING COHORT STATISTICS (LEARNER-BASED SPLIT). LEARNER, KNOWLEDGE-COMPONENT (KC), INTERACTION, AND TEST-EVENT COUNTS ARE POST-PROCESSING TOTALS, NOT RAW PUBLIC-DUMP COUNTS.

DATASET	LEARNERS	KCS	INTERACTIONS	TEST EVENTS
ASSISTMENTS 2012	27,806	265	2.66M	534,150
JUNYI ACADEMY	71,014	1,326	16.2M	3,269,022
XES3G5M	18,066	866	7.95M	1,589,145

B. Splits and seeds

The learner-based split is primary. Users are shuffled with the fold seed; 20% of learners are held out as test, 10% as validation, and the remainder as training. Within a fold, train, validation, and test learners are disjoint. The temporal split is complementary: all interactions are sorted by timestamp; the earliest 70% are training, the next 10% validation, and the latest 20% test. Gate numbers below are not taken from the temporal split. KC-frequency buckets are constructed from the training file only.

Deep models are trained in five runs at seeds 42, 2024, 2025, 2026, and 2027 (fold_0 through fold_4). Seeds 2025 and 2026 share one student partition (fold_2 = fold_3) on all three datasets. These are five training runs and four unique student partitions, not five independent folds. Tables that report mean $\pm$ sd therefore summarize four unique partitions: the two initializations on the duplicated split are averaged first. The gate-robustness table still lists all five training runs across four unique learner partitions.

C. Model settings

Main tables use local IRT, DKT, and SimpleKT implementations, not a pinned pyKT commit (training snapshot commit: NOT RECOVERED). The local SimpleKT class is a two-layer Transformer encoder and is not byte-identical to the official SimpleKT checkpoint [4]. Recovered hyperparameters follow; missing fields are marked NOT RECOVERED and are not imputed.

RECOVERED TRAINING SETTINGS FOR THE MAIN BASELINES. NOT RECOVERED CELLS WERE NOT IMPUTED. THIS LISTING IS NOT A RESULTS TABLE.

SETTING	IRT 1PL	DKT	SIMPLEKT
IMPLEMENTATION	LOCAL IRT	LOCAL DKT	LOCAL SIMPLEKT
VERSION/COMMIT	NOT RECOVERED	NOT RECOVERED	NOT RECOVERED
MAJOR HYPERPARAMETERS	1PL: SIGMOID($\theta-B$); L2 0.01	LSTM EMBED 64, HIDDEN 128	TRANSFORMER 2 LAYERS, 4 HEADS, EMBED 64
OPTIMIZER	SGD	ADAM	ADAM
LEARNING RATE	0.01	1E-3	1E-3
EARLY STOPPING	NONE (10 EPOCHS)	NONE (50 EPOCHS)	NONE (50 EPOCHS)
BATCH SIZE	512	64	64
MAX. SEQUENCE LENGTH	N/A	200	200
SELECTION METRIC	LAST EPOCH	VALIDATION AUC	VALIDATION AUC

On ASSISTments fold 0 only we also score a train-only GKT graph [8] (pyKT GKT; Adam, learning rate 1e-3, batch size 16, maximum length 100, hidden and embedding size 32, dropout 0.5, 20 epochs, patience 4, selection by validation AUC) and a CL4KT protocol adapter [9] (contrastive views on training sequences; not an official CL4KT checkpoint; Adam, learning rate 1e-3, batch size 64, maximum length 100, hidden size 64, 4 heads, 2 layers, dropout 0.2, 20 epochs, patience 6, selection by validation AUC). IRT under learner-based splits has no ability parameter for unseen students, so its AUC is 0.50 by construction; we report it as a base-rate reference, not as a ranking competitor.

D. Train-only frequency strata

For each fold, KC frequency f_{train} is counted on the training file only. Buckets: strict cold-start ($f=0$), very sparse ($0 < f < 20$), sparse ($20 \leq f < 100$), medium ($100 \leq f < 500$), dense ($f \geq 500$). Fig. 1 (top) shows the number of KCs in each stratum. Training-interaction volume is shown separately in the bottom row from summed f_{train} , not inferred from KC counts.




Fig. 1. Distribution of KCs across train-only frequency strata under learner-based and temporal splits. Bottom: verified training-interaction counts (sum of train-only f_{train} on fold 0); volume is not inferred from KC counts. Bottom y-axis is logarithmic.

E. Reliability flags

Occupancy flags follow test-event count N : Reliable (R) $N \geq 1000$; Limited (L) $100 \leq N < 1000$; Insufficient (I) $N < 100$. We call R, L, and I descriptive sample-support flags. They are not inferential guarantees, confidence intervals, or hypothesis tests. Success claims require Limited or Reliable occupancy. Insufficient cells are descriptive only. Very-sparse ASSISTments cells are Insufficient and descriptive only.

F. Calibration

Let $y \in \{0, 1\}$ be next-response correctness and $p \in [0, 1]$ the predicted probability. With $M=15$ equal-width bins, $\text{ECE} = \sum_m (n_m / N) |\text{acc}_m - \text{conf}_m|$. The Brier score is $(1/N) \sum_i (p_i - y_i)^2$. We use the binned decomposition $\text{Brier} = \text{UNC} - \text{RES} + \text{REL}$ [14], with $\text{UNC} = \bar{y}(1 - \bar{y})$, $\text{REL} = (1/N) \sum_m n_m (\text{conf}_m - \text{acc}_m)^2$, and $\text{RES} = (1/N) \sum_m n_m (\text{acc}_m - \bar{y})^2$. Because probabilities are binned, the three components are an empirical approximation and need not sum exactly to the directly computed Brier score. Lower ECE is better calibration; it is not a substitute for AUC.

G. Difficulty coupling

For each KC c we define a training-only difficulty proxy $\text{difficulty}(c) = 1 - \text{mean_train_correctness}(c)$. The mean is taken on the training file only. This is an observational proxy for how often c is answered incorrectly in training; it is not a latent IRT difficulty parameter. The reported Spearman

association is $\rho(\log(1 + f_{\text{train}}), \text{difficulty})$, computed on training-fold KC summaries. The same file yields item support (distinct items per KC), learner exposure (distinct learners per KC), and a curriculum-position proxy (median normalized sequence position). Results report between-KC associations and a within-KC sparsification check; neither identifies a causal frequency effect.

H. Simulated decision gate

For a locked global threshold τ , the system advances (skips remediation) if $p \geq \tau$ and triggers remediation otherwise. Sparse events never receive their own tuned τ . The display threshold is $\tau=0.7$; the grid $\{0.5, 0.6, 0.7, 0.8\}$ is recorded in the artifact and is not a search over sparse error.

Let $A = \{p \geq \tau\}$. The false-advance rate (FAR) is $P(y=0 | p \geq \tau)$: among responses for which the simulated system would advance the learner, the proportion whose observed next response is incorrect. This is not a measurement of latent mastery. It is a next-response decision-error proxy under a simulated threshold gate. $\text{Miss} = P(p \geq \tau | y=0)$. Expected FAR under the model probabilities is $E[1 - p | p \geq \tau]$. $\text{Excess FAR} = \text{FAR} - E[1 - p | p \geq \tau]$. $\Delta\text{FAR} = \text{FAR}_{\text{sparse}} - \text{FAR}_{\text{dense}}$. FAR is a proportion of $N_{\text{advance}} = \text{count}(p \geq \tau)$ events, not of N_{total} ; Miss is a proportion of $N_{\text{incorrect}} = \text{count}(y=0)$. N_{total} alone does not imply the precision of FAR. Where prediction-level exports exist, 95% intervals use the KC-clustered percentile bootstrap already used for this gate ($B=2000$, resample KCs within sparse and within dense). Positive ΔFAR means sparse advances are dirtier than dense

advances. This is a simulated decision error, not an instructional RCT.

IV. RESULT AND DISCUSSION

A. Aggregate Ranking Is Not the Sparse Story

Table 2 shows learner-based AUC. DKT is slightly above SimpleKT on all three datasets. IRT AUC is 0.50 for the reason given above. These numbers would look like a routine KT bake-off. They do not tell a designer whether p is trustworthy on the tail.

TABLE 2. OVERALL LEARNER-BASED AREA UNDER THE ROC CURVE (AUC) AND ACCURACY (ACC) (MEAN±SD OVER FOUR UNIQUE PARTITIONS).

Dataset	Model	AUC	ACC
ASSISTments 2012	IRT	0.5000	0.6973±0.0004
ASSISTments 2012	DKT	0.6979±0.0014	0.7182±0.0014
ASSISTments 2012	SimpleKT	0.6837±0.0025	0.6996±0.0032
Junyi Academy	IRT	0.5000	0.7053±0.0018
Junyi Academy	DKT	0.7320±0.0009	0.7343±0.0015
Junyi Academy	SimpleKT	0.7231±0.0030	0.7274±0.0015
XES3G5M	IRT	0.5000	0.7961±0.0031
XES3G5M	DKT	0.8171±0.0022	0.8327±0.0032
XES3G5M	SimpleKT	0.7557±0.0013	0.8067±0.0037

TABLE 3. SIMPLEKT EVENT-LEVEL EXPECTED CALIBRATION ERROR (ECE) BY TRAIN-ONLY FREQUENCY STRATUM. N IS THE MEAN TEST-EVENT COUNT ACROSS FOUR UNIQUE LEARNER PARTITIONS, NOT A SINGLE-PARTITION COUNT. FLAGS: RELIABLE (R) $N \geq 1000$; LIMITED (L) $100 \leq N < 1000$. JUNYI SPARSE IS EMPTY.

Dataset	Stratum	N	Flag	ECE
ASSISTments 2012	dense	523,971	R	0.1136±0.0066
ASSISTments 2012	medium	5,963	R	0.1541±0.0051
ASSISTments 2012	sparse	415	L	0.2280±0.0197
Junyi Academy	dense	3,232,614	R	0.0792±0.0051
Junyi Academy	medium	3,836	R	0.1073±0.0156
XES3G5M	dense	1,268,696	R	0.1145±0.0011
XES3G5M	medium	12,980	R	0.1114±0.0076
XES3G5M	sparse	2,010	R	0.1248±0.0085

A. A Global Gate Can Turn the ASSISTments Gradient into a Decision Error

Table 4 applies $\tau=0.7$ on ASSISTments fold 0 (seed 42; sparse $N_{\text{total}}=444$, Limited). SimpleKT FAR is 0.196 on dense KCs and 0.268 on sparse KCs ($\Delta\text{FAR}=+0.072$). The FAR denominator is N_{advance} (284,326 dense; 235 sparse), not N_{total} (528,018 and 444); $N_{\text{incorrect}}$ is 158,623 dense

On XES3G5M, sparse AUC is higher than dense AUC (DKT 0.857 vs. 0.817; SimpleKT 0.847 vs. 0.755; Reliable occupancy). Lower frequency is therefore not a universal ranking failure. Junyi’s learner-based sparse bucket is empty under the pre-registered cuts—a protocol outcome, not a missing cell to impute.

B. Calibration Can Move When Ranking Does Not

Table 3 is the calibration punchline. On ASSISTments 2012, SimpleKT ECE increases from 0.1136±0.0066 (dense, Reliable, $N=523,971$) to 0.1541 (medium) to 0.2280±0.0197 (sparse, Limited, $N=415$). DKT shows the same direction (dense ECE 0.060 vs. sparse 0.233). IRT dense ECE is tiny (0.003) but resolution is zero: a constant-like predictor can look “calibrated” while ranking nothing.

On Junyi, only dense and medium strata exist; SimpleKT ECE rises modestly from 0.079 to 0.107. On XES3G5M, SimpleKT ECE is essentially flat (0.114→0.111→0.125 dense/medium/sparse) despite Reliable sparse occupancy ($N=2,010$). ECE-flat is not a license to skip occupancy reporting, and—as the gate results show—it is not the same as a flat miss rate.

and 197 sparse. Expected FAR on sparse advances is 0.050, so Excess FAR is 0.218. DKT shows a similar FAR gap. Train-only GKT shrinks ΔFAR to +0.015 (FAR 0.205→0.220; Excess FAR 0.063 on sparse). The CL4KT adapter is intermediate (0.185→0.240). We do not read this as “use GKT in production” ; we read it as: the ASSISTments calibration gradient is not a property of the dataset alone.

TABLE 4. SIMULATED GATE AT $\tau=0.7$, ASSISTMENTS 2012 FOLD 0. FALSE-ADVANCE RATE (FAR)= $P(Y=0 | p \geq \tau)$ AMONG N_{advance} ADVANCES; MISS AMONG $N_{\text{incorrect}}$ INCORRECT RESPONSES; EXCESS FAR= $\text{FAR}-E[\text{FAR}]$. FAR 95% CIs: KC-CLUSTER PERCENTILE, $B=2000$. NOT LATENT MASTERY; NOT A CLASSROOM TRIAL.

Model	Stratum	N	N_{advance}	$N_{\text{incorrect}}$	FAR [95% CI]	$E[\text{FAR}]$	Excess FAR	Miss
SimpleKT	dense	528,018	284,326	158,623	0.196 [0.186, 0.208]	0.113	0.083	0.352
SimpleKT	sparse	444	235	197	0.268 [0.202, 0.337]	0.050	0.218	0.320
DKT	dense	528,018	315,650	158,623	0.200 [0.190, 0.211]	0.147	0.053	0.398
DKT	sparse	444	243	197	0.296 [0.221, 0.383]	0.057	0.239	0.365
GKT (train-only)	dense	528,018	351,503	158,623	0.205 [0.194, 0.217]	0.163	0.042	0.455
GKT (train-only)	sparse	444	209	197	0.220 [0.149, 0.295]	0.157	0.063	0.234
CL4KT (adapter)	dense	528,018	307,479	158,623	0.185 [0.176, 0.194]	0.175	0.010	0.359
CL4KT (adapter)	sparse	444	200	197	0.240 [0.159, 0.330]	0.116	0.124	0.244

Seed-42 ΔFAR 95% CI (KC-cluster, $B=2000$): SimpleKT [0.006, 0.138] (locked C2); DKT [0.019, 0.175]; GKT [-0.054, 0.092]; CL4KT [-0.018, 0.142].

Table 5 checks whether SimpleKT ΔFAR is a one-fold

accident. It is positive on all five training runs across four unique learner partitions (mean 0.047, sd 0.033). Mean sparse denominators are $N_{\text{total}}=413$, $N_{\text{advance}}=227$, $N_{\text{incorrect}}=155$. A KC-clustered bootstrap on seed 42

yields a 95% interval [0.006, 0.138] for ΔFAR —excluding 0, but wide, as Limited occupancy and $N_{\text{advance}}=235$ require. DKT ΔFAR is positive on only three of five seeds and is not treated as a five-run finding. On XES3G5M, SimpleKT ΔFAR is negative on all five seeds, but ΔMiss is positive on all five (mean +0.112): among actual incorrect answers, the

system still advances more often on sparse than on dense KCs. A flat ECE therefore does not imply a flat miss rate.

Coincidence of digits: SimpleKT dense $E[\text{FAR}]$ at $\tau=0.7$ on seed 42 is 0.113, close to the four-partition dense ECE 0.114. They are different quantities.

TABLE 5. GATE ROBUSTNESS AT $\tau=0.7$ ON ASSISTMENTS 2012. FIVE TRAINING RUNS (SEEDS 42, 2024, 2025, 2026, 2027) ACROSS FOUR UNIQUE LEARNER PARTITIONS (2025 AND 2026 SHARE A SPLIT). $\Delta\text{FAR}=\text{FAR}_{\text{SPARSE}}-\text{FAR}_{\text{DENSE}}$. MEAN N , N_{ADVANCE} , AND $N_{\text{INCORRECT}}$ ARE SPARSE-STRATUM DENOMINATORS. SIMPLEKT: $\Delta\text{FAR}>0$ ON 5/5 TRAINING RUNS AND ON ALL FOUR UNIQUE PARTITIONS. DKT: $\Delta\text{FAR}>0$ ON 3/5 TRAINING RUNS (THE SHARED PARTITION IS MIXED IN SIGN) AND IS NOT A FIVE-RUN FINDING. GKT/CL4KT REMAIN SEED 42 ONLY.

Model	Mean ΔFAR	SD	Runs $\Delta\text{FAR}>0$	Mean N	Mean N_{advance}	Mean $N_{\text{incorrect}}$
SimpleKT	0.047	0.033	5/5 runs (4 partitions)	413	227	155
DKT	0.033	0.048	3/5 runs	413	226	155

A. What a Designer Should Take Away

If a learning platform uses KT probabilities only to sort students, Table 2 may suffice. If it uses p to skip practice or declare mastery, Table 2 is the wrong dashboard. The ASSISTments SimpleKT cell is the cautionary case: competitive AUC, Limited-but-estimable sparse occupancy (about 19% of KCs fall below the sparse threshold on fold 0), a monotonic ECE rise, and a gate that advances more dirty sparse attempts. The XES3G5M cell is the opposite caution: plenty of sparse mass and Reliable occupancy, yet SimpleKT ECE does not degrade—while miss rates still can. Junyi shows that the protocol can refuse a sparse claim when the bucket is empty.

On the three datasets studied here, three empirical conditions track when a sparse-calibration contrast is estimable rather than empty (Table 6): (1) a non-empty low-frequency tail under the split actually used; (2) enough sparse test events for at least a Limited-flag ECE; (3) frequency–difficulty coupling that is not inverted. ASSISTments 2012 meets all three. That is a diagnostic pre-condition on these logs, not a universal law and not a claim that training frequency causes miscalibration.

TABLE 6. EMPIRICAL OBSERVATIONS, ON THESE THREE DATASETS, OF WHEN A SPARSE-CALIBRATION CONTRAST IS ESTIMABLE. THESE ARE NOT UNIVERSAL LAWS. SPARSE MASS: SHARE OF KCs WITH TRAIN-ONLY FREQUENCY $f_{\text{TRAIN}}<100$.

Condition	ASSISTments	Junyi	XES3G5M
Sparse mass	18.9% of KCs	0%	22.5% of KCs
Sparse N	415 (L)	empty	2,010 (R)
Difficulty coupling	$\rho=-0.227$ (sparse harder)	$\rho=-0.416$	$\rho=+0.087$ (weak, inverted)
SimpleKT ECE	0.114→0.228	dense→medium only	flat 0.114→ 0.125

B. Explanatory Analysis of Dataset-Dependent Calibration

Table 6 records three observational pre-conditions under which a sparse-calibration contrast is estimable on these logs. The same training-only file also yields item support, learner exposure, and a curriculum-position proxy. This subsection reports those descriptors and two distinct estimands. Neither is a causal claim.

Sparse mass (share of KCs with $f_{\text{train}}<100$) is 18.9% on ASSISTments 2012, 0% on Junyi Academy, and 22.5% on

XES3G5M. Sparse test-event support is Limited ($N=415$), empty, and Reliable ($N=2,010$). Frequency–difficulty coupling (Spearman ρ of $\log(1+f_{\text{train}})$ with $1-\text{mean train correctness}$) is $\rho=-0.227$ and $\rho=-0.416$ (rarer KCs harder) versus a weak opposite-signed $\rho=+0.087$ on XES3G5M. Item support: ASSISTments is item-rich (median 44.5 items/KC; dense 205 vs sparse 1); Junyi is nearly uniform (median 18, IQR 5); XES3G5M is thin (median 3; dense 9 vs sparse 1). Learner exposure is the count of distinct training learners per KC. Curriculum-position coupling is $\rho=-0.308$, -0.324 , and -0.125 .

A test-event-weighted KC-level regression of SimpleKT expected calibration error (ECE) on those five training-only covariates ($n=478$, 2,645, and 1,263 KCs) is a between-KC association: it compares different concepts. Standardized $\log(1+f_{\text{train}})$ coefficients are -0.079 [-0.097 , -0.061] on ASSISTments, -0.010 [-0.014 , -0.007] on Junyi, and -0.117 [-0.171 , -0.063] on XES3G5M in the weighted fit. Difficulty is independently associated on ASSISTments and Junyi. Learner exposure is independently associated only on ASSISTments ($+0.019$ [0.010 , 0.027]). These signs do not test a frequency dose on a fixed KC.

Within-KC controlled sparsification holds KC identity, the test set, labels, and all other KCs’ training rows fixed, and reduces training rows for 30 originally dense KCs ($f_{\text{train}}\geq 500$; pre-registered selection, seed 42, fold 0) to 500 or 50 rows. Table 7 reports key paired ΔECE (reduced–full), with KC-bootstrap 95% CIs. Reducing training evidence for the same KC does not universally worsen calibration: several CIs lie below 0 or include 0. The observational ASSISTments SimpleKT dense-to-sparse ECE gradient (0.114→0.228) is not reproduced by sparsifying originally dense ASSISTments KCs (SimpleKT, 50 rows: $+0.002$ [-0.021 , $+0.025$]). Frequency alone is therefore not a universal causal explanation. Junyi SimpleKT at 50 rows does show a large positive ΔECE ; that cell shows a within-KC increase is possible, not that it is a law.

Table 7. Within-KC controlled sparsification (seed 42; 30 originally dense KCs per dataset). Delta ECE = reduced–full; positive means worse calibration after keeping 500 or 50 training rows for the same KC. 95% CIs bootstrap over KCs. Not a causal law for real-world sparsity.

Dataset	Model	Reduction	Delta ECE	95% CI	Interpretation
ASSISTments 2012	DKT	500 rows	−0.047	[−0.060, −0.033]	ECE lower
ASSISTments 2012	SimpleKT	50 rows	+0.002	[−0.021, +0.025]	CI includes 0
Junyi Academy	DKT	500 rows	−0.021	[−0.041, −0.001]	ECE lower
Junyi Academy	SimpleKT	50 rows	+0.135	[+0.110, +0.161]	ECE higher
XES3G5M	DKT	500 rows	−0.008	[−0.019, +0.004]	CI includes 0

A. What This Paper Does Not Show

The gate is not a classroom policy and not an A/B test of remediation. GKT and the CL4KT adapter are single-fold, ASSISTments-only instantiations. Temporal results are a single corrected cutoff (seed 42), not a multi-cutoff variance estimate. IRT is not a fair ranking baseline on unseen learners. We did not tune τ on sparse events, and we do not recommend doing so in deployment without a separate validation design.

B. Implications for Information and Education Technology

IJET’s audience includes operators of educational platforms, not only KT researchers. For that audience, the operational recommendation is narrow: (i) log train-only KC frequency with every prediction; (ii) report ECE and occupancy on dense vs. sparse slices before enabling a global mastery threshold; (iii) if a threshold must be global, simulate FAR and Miss on the sparse slice at the same τ used in the product; (iv) do not treat a population AUC win, or a graph/contrastive architecture, as automatically safer on the tail.

V. CONCLUSION

KT models that look adequate on AUC can still be poorly calibrated on rarely practiced skills, and a global mastery threshold can then produce a higher false-advance rate on those skills. That pattern appears for SimpleKT on ASSISTments 2012 (ECE 0.114→0.228, Limited N=415; seed-42 FAR 0.196→0.268; Δ FAR>0 on 5/5 training runs across four unique learner partitions). It does not appear as a universal law: Junyi has no learner-based sparse stratum, and XES3G5M SimpleKT ECE is flat. Sparse-concept diagnostics are therefore conditionally important. The simulation is a decision-error check for educational-technology systems, not a claim of a new KT model and not a classroom intervention.

CONFLICT OF INTEREST

The authors declare no conflict of interest.

AUTHOR CONTRIBUTIONS

K.-T.N. led the diagnostic design, experiments, and manuscript. T.D.M. and D.N.T. contributed data processing and baseline runs. C.T.N. contributed methodological review. V.-H.N. supervised the study and revised the

manuscript. All authors approved the final version.

ACKNOWLEDGMENT

Anonymized code and prediction exports are available for peer review; a public repository will be released upon acceptance. This study uses de-identified public learner logs (ASSISTments 2012, Junyi Academy, XES3G5M) and is not a classroom trial.

REFERENCES

- [1] A. T. Corbett and J. R. Anderson, “Knowledge Tracing: Modeling the acquisition of procedural knowledge,” *User Modeling and User-Adapted Interaction*, vol. 4, no. 4, pp. 253–278, 1994.
- [2] C. Piech, J. Bassen, J. Huang, S. Ganguli, M. Sahami, L. J. Guibas, and J. Sohl-Dickstein, “Deep Knowledge Tracing,” in *Advances in Neural Information Processing Systems*, 2015, pp. 505–513.
- [3] G. Abdelrahman, Q. Wang, and B. Nunes, “Knowledge Tracing: A survey,” *ACM Computing Surveys*, vol. 55, no. 11, pp. 1–37, 2023.
- [4] Z. Liu, Q. Liu, J. Chen, S. Huang, and W. Luo, “SimpleKT: A simple but tough-to-beat baseline for Knowledge Tracing,” in *The Eleventh International Conference on Learning Representations*, 2023.
- [5] Z. Liu, Q. Liu, J. Chen, S. Huang, J. Tang, and W. Luo, “pyKT: A python library to benchmark deep learning based Knowledge Tracing models,” in *Thirty-sixth Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2022.
- [6] G. Rasch, *Probabilistic Models for Some Intelligence and Attainment Tests*. Danish Institute for Educational Research, 1960.
- [7] A. Ghosh, N. Heffernan, and A. S. Lan, “Context-aware attentive Knowledge Tracing,” in *Proc. ACM SIGKDD Conf. Knowledge Discovery and Data Mining*, 2020, pp. 2330–2339.
- [8] H. Nakagawa, Y. Iwasawa, and Y. Matsuo, “Graph-based Knowledge Tracing: Modeling student proficiency using graph neural network,” in *IEEE/WIC/ACM International Conference on Web Intelligence*, 2019, pp. 156–163.
- [9] W. Lee, J. Chun, Y. Lee, K. Park, and S. Park, “Contrastive learning for Knowledge Tracing,” in *Proc. ACM Web Conference*, 2022, pp. 2330–2338.
- [10] R. Pelánek, “Metrics for evaluation of student models,” *Journal of Educational Data Mining*, vol. 7, no. 2, pp. 1–19, 2015.
- [11] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, “On calibration of modern neural networks,” in *Proc. International Conference on Machine Learning*, 2017, pp. 1321–1330.
- [12] M. P. Naeini, G. Cooper, and M. Hauskrecht, “Obtaining well calibrated probabilities using Bayesian binning,” in *Proc. AAAI Conference on Artificial Intelligence*, vol. 29, no. 1, 2015.
- [13] G. W. Brier, “Verification of forecasts expressed in terms of probability,” *Monthly Weather Review*, vol. 78, no. 1, pp. 1–3, 1950.
- [14] M. H. DeGroot and S. E. Fienberg, “The comparison and evaluation of forecasters,” *The Statistician*, vol. 32, no. 1/2, pp. 12–22, 1983.
- [15] X. Yan, C. Tang, and A. Shimada, “Recovering stranded discrimination in Knowledge Tracing: Per-item bias correction via empirical-Bayes shrinkage,” *arXiv:2606.14123*, 2026.
- [16] S. Huang, Z. Liu, X. Zhao, W. Luo, and J. Weng, “Towards robust Knowledge Tracing models via k-sparse attention,” in *Proc. 46th International ACM SIGIR Conference*, 2023, pp. 2441–2445.
- [17] I. Bhattacharjee and C. Wayllace, “Cold start problem: An experimental study of Knowledge Tracing models with new students,”

in Artificial Intelligence in Education (AIED 2025), LNCS vol. 15880, Springer, 2025, pp. 425–432.

- [18] ASSISTmentsData, “ASSISTments 2012–2013 school data with affect,” ASSISTments Public Datasets, 2012. [Online]. Available: <https://sites.google.com/site/assistmentsdata/datasets/2012-13-school-data-with-affect>
- [19] Junyi Academy, “Junyi Academy online learning activity dataset,” Kaggle, 2019. [Online]. Available: <https://www.kaggle.com/datasets/junyiacademy/learning-activity-public-dataset-by-junyi-academy>

- [20] Z. Liu, Q. Liu, T. Guo, J. Chen, S. Huang, X. Zhao, J. Tang, W. Luo, and J. Weng, “XES3G5M: A Knowledge Tracing benchmark dataset with auxiliary information,” in Thirty-seventh Conference on Neural Information Processing Systems Datasets and Benchmarks Track, 2023.

Copyright © 2026 by the authors. This is an open access article distributed under the Creative Commons Attribution License which permits unrestricted use, distribution, and reproduction in any medium, provided the original work is properly cited ([CC BY 4.0](https://creativecommons.org/licenses/by/4.0/)).